


MIG (Multi-Instance GPU) on Wulver

Configuration • Usage • Best Practices


Agenda

- What is MIG?
- Why MIG on Wulver?
- MIG Configuration Example
- Submitting Jobs (srun & sbatch)
- New Billing Model
- Summary & Q&A

Wulver Specifications




GPU Allocated vs used



GPUs were reserved ~85–95% of the time, while average compute utilization was ~25–50%

GPU Mem Usage



GPUs were **heavily** allocated but **lightly utilized**

GPU Utilization on Wulver

- GPUs were reserved **85–95% of the time**, but:
 - Average compute utilization was only **25–50%**.
 - Memory utilization was mostly **below 20%**.
- This mismatch caused:
 - Long queue times
 - Wasted capacity
 - Full A100 GPUs held by small jobs

What is MIG?

- Partitions an A100 GPU into multiple isolated GPU instances
- Each instance has dedicated memory, SMs(Streaming Multiprocessor), cache, bandwidth
- Appears as a separate GPU to CUDA, PyTorch, TensorFlow
- Up to 7 MIG GPU instances per A100


MIG Profiles on A100 MIG

Config	GPC Slice #0	GPC Slice #1	GPC Slice #2	GPC Slice #3	GPC Slice #4	GPC Slice #5	GPC Slice #6	OFA	NVDEC	NVJPG	P2P	GPU Direct RDMA
1					7			1	5	1	No	
2			4			3		0	2+2	0	No	
3			4			2	1	0	2+1+0	0	No	
4			4		1	1	1	0	2+0+0+0	0	No	
5		3			3			0	2+2	0	No	
6		3			2	1		0	2+1+0	0	No	
7		3		1	1	1		0	2+0+0+0	0	No	
8	2		2			3		0	1+1+2	0	No	
9	2		1	1		3		0	1+0+0+2	0	No	
10	1	1		2		3		0	0+0+1+2	0	No	
11	1	1	1	1		3		0	0+0+0+0+2	0	No	
12	2		2		2	1		0	1+1+1+0	0	No	
13	2		1	1	2	1		0	1+0+0+1+0	0	No	
14	1	1		2	2	1		0	0+0+1+1+0	0	No	
15	2		1	1	1	1		0	1+0+0+0+0	0	No	
16	1	1		2	1	1	1	0	0+0+1+0+0+0	0	No	
17	1	1	1	1	2	1		0	0+0+0+0+1+0	0	No	
18	1	1	1	1	1	2		0	0+0+0+0+0+1	0	No	
19	1	1	1	1	1	1	1	0	0+0+0+0+0+0+0	0	No	


Supported
MemBW
proportional
to size of the
instance

Reference: <https://docs.nvidia.com/datacenter/tesla/mig-user-guide/index.html#a100-mig-profiles>

MIG Configuration Example



MIG Configuration Example –Cont.



Source: <https://developer.nvidia.com/techdemos/video/disc03>

Why MIG on Wulver?

- Reduce queue times: smaller slices fit into schedule easier.
- Improve overall cluster utilization.
- Fair, predictable performance via hardware isolation.
- Lower SU cost for small/medium jobs compared to full A100.

MIG Configuration on Wulver

```
$ nvidia-smi -L
GPU 0: NVIDIA A100-SXM4-80GB (UUID: GPU-7e92a539-19c5-2948-67d5-a3b055a816ed)
  MIG 3g.40gb    Device  0: (UUID: MIG-0b586768-e667-5a5c-8298-7cfceacc60fe)
  MIG 2g.20gb    Device  1: (UUID: MIG-0df8d316-2b51-55d5-9d6f-c78f4edc27f4)
  MIG 1g.10gb    Device  2: (UUID: MIG-379a24be-c1b3-5d57-959a-814f18460157)
  MIG 1g.10gb    Device  3: (UUID: MIG-fc813cd3-631f-5a59-9ead-3ac7ed3215db)
GPU 1: NVIDIA A100-SXM4-80GB (UUID: GPU-377e6263-247c-2736-6b78-6439f6164e94)
  MIG 3g.40gb    Device  0: (UUID: MIG-079d02a7-0892-5ea1-88f5-ca22b52ee4f3)
  MIG 2g.20gb    Device  1: (UUID: MIG-978fe824-62ed-5bc7-8f96-c00e09dfec7)
  MIG 1g.10gb    Device  2: (UUID: MIG-dfdb3d5e-c3f6-5189-b355-7a65c8a8368e)
  MIG 1g.10gb    Device  3: (UUID: MIG-ed0b10c7-3891-5f2e-a02c-ece69c8ce419)
GPU 2: NVIDIA A100-SXM4-80GB (UUID: GPU-c479faea-dea9-fc74-8029-0c5bcd509ec7)
  MIG 3g.40gb    Device  0: (UUID: MIG-ec7618a0-0e1a-5919-a88e-5c9ca61cebad)
  MIG 2g.20gb    Device  1: (UUID: MIG-9ec817d8-06b6-5ac0-94fc-9ae1c233d506)
  MIG 1g.10gb    Device  2: (UUID: MIG-0f10402e-629d-5f98-a838-0b3798fc842)
  MIG 1g.10gb    Device  3: (UUID: MIG-97b3e81f-d23e-50ee-b432-c1938638f6f4)
GPU 3: NVIDIA A100-SXM4-80GB (UUID: GPU-3d974751-2093-6cf6-111a-3a8c95f68a97)
  MIG 3g.40gb    Device  0: (UUID: MIG-d0128346-7241-502e-99bf-e0e73e96be4d)
  MIG 2g.20gb    Device  1: (UUID: MIG-775a5688-9a5c-54c0-b481-c514414d6149)
  MIG 1g.10gb    Device  2: (UUID: MIG-7a0419b4-6c67-5ba0-aa66-b68778fc8bdc)
  MIG 1g.10gb    Device  3: (UUID: MIG-1d33f90a-1571-5770-88e9-6505a4a6ddf9)
```

MIG Profiles on Wulver

- MIG
 - a100_10g → ~10 GB slice
 - a100_20g → ~20 GB slice
 - a100_40g → ~40 GB slice
- full GPU
 - a100 (--gres=gpu:a100:1) (80 GB)

How MIG Appears in SLURM

- `--gres=gpu:a100_10g:1`
- `--gres=gpu:a100_20g:1`
- `--gres=gpu:a100_40g:1`
- `--gres=gpu:a100:1` (full GPU)
- Each MIG slice is treated as a unique GPU device

GPU MIG	Slurm Directive
10G MIG	<code>--gres=gpu:a100_10g:1</code>
20G MIG	<code>--gres=gpu:a100_20g:1</code>
40G MIG	<code>--gres=gpu:a100_40g:1</code>

Submitting Jobs (Batch)

```
#!/bin/bash -l
#SBATCH --job-name=mig_test
#SBATCH --output=%x.%j.out
#SBATCH --error=%x.%j.err
#SBATCH --partition=gpu
#SBATCH --qos=standard
#SBATCH --account=$PI
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --mem-per-cpu=4000M
#SBATCH --gres=gpu:a100_40g:1
#SBATCH --time=02:00:00

module load Miniforge3
conda activate torch-cuda
srun python torch_tensor.py
```

Submit Jobs (Interactive)

```
$srun --partition=gpu \
    --account=$PI_ucid \
    --qos=standard \
    --gres=gpu:a100_10g:1 \
    --time=00:59:00 \
    --pty bash
```

MIG in Open OnDemand

Servers

- Jupyter
- MATLAB Server
- Jupyter-matlab-proxy
- VS Code
- RStudio

Conda Environment to be Activated

tf

- Select the name of environment to be activated. This field is mandatory if you choose Conda environment.
- Your environment must have Jupyter Notebook. Check Jupyter Documentation for details.

Working directory

Select Path

Select your project directory; defaults to \$HOME

Enter any prerequisite commands if required

Optional (example: `module load` command or `source /path/to/bin/activate` if using Python virtual env)

Account

walsh

Partition

gpu

QOS

standard

Number of hours (max 72)

1

Maximum wall time requested

Number of cores (max 128)

1

Number of cores on node

Select number of GPUs or MIG Device

✓ 1

2

3

4

MIG 40GB

MIG 20 GB

MIG 10 GB

* For example, if 4 cores and 128GB of memory are requested,


Key Factors for Choosing a MIG Slice

- **Compute requirements** → CUDA cores, Tensor cores, bandwidth
- **Memory needs** → match slice memory to workload footprint
- **Concurrency** → number of simultaneous jobs / services
- **Latency sensitivity** → real-time or interactive workloads
- **Workload profile** → compute-bound vs memory-bound

When to Use Which MIG Slice?

- 10g → inference, debugging, preprocessing
- 20g → medium training, light ML workloads
- 40g → larger models, heavy DL
- 80g (full GPU) → GPU memory-intensive or high-performance workloads

Tools to Profile GPU jobs



Best Practices

- Use smallest MIG slice needed
- Avoid over-allocating CPU or RAM
- Monitor GPU usage with nvidia-smi or other profiling tools
- Track job resource usage for reproducibility

Service Unit (SU) Calculation

SU = MAX(#CPUs, Memory(in GB)/4) + 16 × (GPU Memory requested / 80GB)

- CPU/Memory term: Accounts for CPU cores and RAM usage
- GPU term: Scales with GPU memory requested
- Applies to both MIG slices and full A100 GPUs
- HPC Service Units resource
page: https://hpc.njit.edu/Running_jobs/service-units/

How SU Charges Are Applied

- CPU and memory usage are always included in SU calculations
- Full A100 GPU (80 GB) → billed as 16 SU/hr
- MIG slices → billed in fractions, proportional to GPU memory share

Billing Examples (NJIT, 4 CPUs)

Profile	GPU Memory	CPU Alloc	SU/hr
a100_10g	10 GB	4	6
a100_20g	20 GB	4	8
a100_40g	40 GB	4	12
Full A100	80 GB	4	20

Example 1: 4 CPUs + full A100 GPU (80 GB)

$$SU = \text{MAX}(4, 4*4G/4) + 16 \times (80/80) = 20 \text{ SU/hr}$$

Example 2: 4 CPUs + MIG slice (10 GB)

$$SU = \text{MAX}(4, 4*4G/4) + 16 \times (10/80) = 6 \text{ SU/hr}$$

Useful Links

- MIG Documentation: <https://hpc.njit.edu/MIG/>
- Service Units:
https://hpc.njit.edu/Running_jobs/service-units/
- Events & Workshops:
https://hpc.njit.edu/HPC_Events_and_Workshops/

Summary & Q&A

- MIG enables right-sizing:
a100_10g / 20g / 40g / full GPU.
- SU charges are now calculated more precisely,
incorporating full details for CPU, memory, and
GPU requests.
- MIG available in different partition
- Questions?